

Business Case: OLA - Ensemble Learning

Problem Statement

Ola faces a major challenge with high driver churn, causing costly recruitment and morale issues. Drivers often leave suddenly or switch to competitors like Uber. The company is hiring widely, including those without cars, which increases acquisition costs. Retaining drivers is more cost-effective than acquiring new ones. The task is to predict driver attrition using monthly data from 2019 and 2020 based on driver attributes.

Project Overview

This project focuses on building a machine learning model to predict **employee churn** — i.e., whether an employee is likely to leave the organization. Reducing churn is critical for minimizing hiring costs, retaining knowledge, and maintaining team performance.

Business Objective

- Identify employees at **risk of leaving**
- Take **proactive actions** (like incentives, promotions, or support)
- Improve **employee retention** and reduce **attrition-related losses**

What This Project Includes

- **Exploratory Data Analysis (EDA)** to understand key patterns and trends
- **Data Preprocessing:** Handling missing values, encoding, KNN and feature engineering
- **Feature Importance Analysis** to find what drives churn
- **Model Building** using ensemble models (Random Forest / XGBoost / DT)
- **Model Tuning** using GridSearchCV and cross-validation
- **Model Evaluation** using metrics like accuracy, precision, recall, F1-score, and confusion matrix

Outcome

- Final model achieved an accuracy of **80%**, with strong recall and F1-score for detecting potential churners.
- **Total Business Value, Income, and Quarterly Rating** were identified as the top drivers of churn.

Jupyter Notebook Analysis

This report presents a comprehensive analysis conducted using Jupyter Notebook. For complete details, including code and visualizations, please refer to the full analysis link provided at the **end of the report**.

1. Data Importing

1.1 Importing the dataset

Code:

```
df=pd.read_csv("Downloads/ola_driver_scaler.csv")
df.head()
```

Output:

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value	Quarterly Rating
0	0	01/01/19	1	28.0	0.0	C23	2	57387	24/12/18	NaN	1	1	2381060	2
1	1	02/01/19	1	28.0	0.0	C23	2	57387	24/12/18	NaN	1	1	-665480	2
2	2	03/01/19	1	28.0	0.0	C23	2	57387	24/12/18	03/11/19	1	1	0	2
3	3	11/01/20	2	31.0	0.0	C7	2	67016	11/06/20	NaN	2	2	0	1
4	4	12/01/20	2	31.0	0.0	C7	2	67016	11/06/20	NaN	2	2	0	1

1.2 Dropping Unnecessary Column

Code:

```
df.drop(columns=['Unnamed: 0'], inplace=True)
```

1.3 Exploring the dataset

Code:

```
df.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   MMM-YY              19104 non-null  object
1   Driver_ID           19104 non-null  int64
2   Age                 19043 non-null  float64
3   Gender              19052 non-null  float64
4   City                19104 non-null  object
5   Education_Level     19104 non-null  int64
6   Income              19104 non-null  int64
7   Dateofjoining       19104 non-null  object
8   LastWorkingDate     1616 non-null   object
9   Joining Designation 19104 non-null  int64
10  Grade               19104 non-null  int64
11  Total Business Value 19104 non-null  int64
12  Quarterly Rating     19104 non-null  int64
dtypes: float64(2), int64(7), object(4)
memory usage: 1.9+ MB
```

1.4 Datetime Conversion for Date Columns

Code:

```
df['MMM-YY'] = pd.to_datetime(df['MMM-YY'])
```

```
##Converting 'Dateofjoining' feature to datetime type
df['Dateofjoining'] = pd.to_datetime(df['Dateofjoining'])
```

```
##Converting 'LastWorkingDate' feature to datetime type
df['LastWorkingDate'] = pd.to_datetime(df['LastWorkingDate'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MMM-YY                19104 non-null  datetime64[ns]
1   Driver_ID             19104 non-null  int64
2   Age                   19043 non-null  float64
3   Gender                19052 non-null  float64
4   City                  19104 non-null  object
5   Education_Level       19104 non-null  int64
6   Income                19104 non-null  int64
7   Dateofjoining         19104 non-null  datetime64[ns]
8   LastWorkingDate       1616 non-null   datetime64[ns]
9   Joining Designation   19104 non-null  int64
10  Grade                 19104 non-null  int64
11  Total Business Value  19104 non-null  int64
12  Quarterly Rating      19104 non-null  int64
dtypes: datetime64[ns](3), float64(2), int64(7), object(1)
memory usage: 1.9+ MB
```

2. Imputation of missing data

2.1 Checking for missing values

```
df.isnull().sum()
```

```
MMM-YY                0
Driver_ID             0
Age                   61
Gender                52
City                  0
Education_Level       0
Income                0
Dateofjoining         0
LastWorkingDate       17488
Joining Designation   0
Grade                 0
Total Business Value  0
Quarterly Rating      0
dtype: int64
```

Inisghts:

LastWorkingDate has 17,488 missing values, likely indicating active drivers.

This can be used to create an attrition flag (1 for left, 0 for active).

Age has 61 missing values and may need imputation using median or similar group.

Gender has 52 missing values and can be filled with mode or a separate "Unknown" category.

All other features have no missing values and are ready for analysis.

2.2 Understanding Categorical Feature Distribution (Gender & Education Level)

```
_data['Gender'].value_counts()
```

```
0.0    11074
1.0     7978
Name: Gender, dtype: int64
```

```
_data['Education_Level'].value_counts()
```

```
1     6864
2     6327
0     5913
Name: Education_Level, dtype: int64
```

Gender Insights:

1. There are **11,074 drivers** labeled as **0.0** and **7,978** as **1.0**.
2. Assuming 0 = Male and 1 = Female, the driver workforce is **male-dominated** (~58% male, ~42% female).
3. This indicates a **gender imbalance** in the dataset.

Education Level Insights:

1. Education levels are categorized as 0, 1, and 2, with **6,864 drivers** at level **1**, **6,327** at level **2**, and **5,913** at level **0**.
2. The distribution is fairly balanced but slightly skewed towards level 1.

2.3 KNN Imputation

```
#knn Imputation
```

```
df_nums=df.select_dtypes(np.number)
```

```
df_nums.head()
```

	Driver_ID	Age	Gender	Education_Level	Income	Joining	Designation	Grade	Total Business Value	Quarterly Rating
0	1	28.0	0.0	2	57387		1	1	2381060	2
1	1	28.0	0.0	2	57387		1	1	-665480	2
2	1	28.0	0.0	2	57387		1	1	0	2
3	2	31.0	0.0	2	67016		2	2	0	1
4	2	31.0	0.0	2	67016		2	2	0	1

```
] : df_nums.isnull().sum()
```

```
] : Driver_ID      0
    Age           61
    Gender        52
    Education_Level 0
    Income         0
    Joining Designation 0
    Grade          0
    Total Business Value 0
    Quarterly Rating 0
    dtype: int64
```

```
df_nums.drop(columns='Driver_ID',inplace=True)
columns=df_nums.columns
```

```
from sklearn.impute import KNNImputer
imputer = KNNImputer(n_neighbors=5, weights='uniform', metric='nan_euclidean',)
imputer.fit(df_nums)
# transform the dataset
df_new = imputer.transform(df_nums)
```

```
df_new=pd.DataFrame(df_new)
```

```
df_new.head()
```

	0	1	2	3	4	5	6	7
0	28.0	0.0	2.0	57387.0	1.0	1.0	2381060.0	2.0
1	28.0	0.0	2.0	57387.0	1.0	1.0	-665480.0	2.0
2	28.0	0.0	2.0	57387.0	1.0	1.0	0.0	2.0
3	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0
4	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0

Insights:

KNN Imputation fills missing values based on the values of the 5 nearest samples, making it **context-aware** and more reliable than mean/median imputation.

The data appears to have been normalized or scaled before imputation, which is good practice for KNN-based methods.

The `df_new.head()` preview confirms that missing values were filled successfully.

It's important to **reassign column names** after imputation (`df_new = pd.DataFrame(df_new, columns=columns)`) for interpretability.

```
: df_new.isnull().sum()
```

```
: Age                0
  Gender              0
  Education_Level     0
  Income              0
  Joining Designation 0
  Grade               0
  Total Business Value 0
  Quarterly Rating    0
  dtype: int64
```

2.4 Reconstructing Final Dataset after KNN Imputation

Code and Output:

```
#Getting Remaning coloumns back  
remaining_columns=list(set(df.columns).difference(set(columns)))
```

```
data=pd.concat([df_new, df[remaining_columns]],axis=1)
```

```
data.head()
```

	Age	Gender	Education_Level	Income	Joining Designation	Grade	Total Business Value	Quarterly Rating	City	MMM-YY	Dateofjoining	LastWorkingDate	Driver_ID
0	28.0	0.0	2.0	57387.0	1.0	1.0	2381060.0	2.0	C23	2019-01-01	2018-12-24	NaT	1
1	28.0	0.0	2.0	57387.0	1.0	1.0	-665480.0	2.0	C23	2019-02-01	2018-12-24	NaT	1
2	28.0	0.0	2.0	57387.0	1.0	1.0	0.0	2.0	C23	2019-03-01	2018-12-24	2019-03-11	1
3	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0	C7	2020-11-01	2020-11-06	NaT	2
4	31.0	0.0	2.0	67016.0	2.0	2.0	0.0	1.0	C7	2020-12-01	2020-11-06	NaT	2

```
df1=pd.DataFrame()
```

```
df1['Driver_ID']=data['Driver_ID'].unique()
```

Insights:

1. The dataset is now fully restructured and contains **no missing values** in the numeric columns after imputation.
2. **LastWorkingDate** still has **NaT values**, which could indicate currently active employees (not necessarily an error).
3. The **negative 'Total Business Value'** (e.g., -665480.0) is suspicious and may need further validation or correction.
4. **City and date columns** were preserved correctly without being affected by imputation.
5. The dataset is now ready for **further analysis or modeling**, with consistent structure and cleaned values.

3. Feature Engineering

3.1 Driver-Level Aggregation of Features from Raw Records

```
#Aggregation at Driver Level

df1['Age'] = list(df.groupby('Driver_ID',axis=0).max('MMM-YY')['Age'])
df1['Gender'] = list(df.groupby('Driver_ID').agg({'Gender':'last'})['Gender'])
df1['City'] = list(df.groupby('Driver_ID').agg({'City':'last'})['City'])
df1['Education'] = list(df.groupby('Driver_ID').agg({'Education_Level':'last'})['Education_Level'])
df1['Income'] = list(df.groupby('Driver_ID').agg({'Income':'last'})['Income'])
df1['Joining_Designation'] = list(df.groupby('Driver_ID').agg({'Joining Designation':'last'})['Joining Designation'])
df1['Grade'] = list(df.groupby('Driver_ID').agg({'Grade':'last'})['Grade'])
df1['Total_Business_Value'] = list(df.groupby('Driver_ID',axis=0).sum('Total Business Value')['Total Business Value'])
df1['Last_Quarterly_Rating'] = list(df.groupby('Driver_ID').agg({'Quarterly Rating':'last'})['Quarterly Rating'])
```

Insights:

1. The data is now **summarized per driver**, making it ideal for modeling or performance comparison.
2. Using .last() ensures the **latest known values** are picked for each driver's profile.
3. **Total Business Value is accumulated** across time, providing a metric of driver contribution.
4. Useful for building KPIs like **income vs. performance**, or analyzing **retention based on recent ratings**.
5. Can now proceed to **segmentation or predictive modeling** using this driver-level dataset.

3.2 Tracking Quarterly Rating Improvements of Drivers

Code:

```
#Quarterly rating at the beginning
qrf = df.groupby('Driver_ID').agg({'Quarterly Rating':'first'})

#Quarterly rating at the end
qrl = df.groupby('Driver_ID').agg({'Quarterly Rating':'last'})

#The dataset which has the employee ids and a boolean value which tells if the rating has increased
qr = (qrl['Quarterly Rating']>qrf['Quarterly Rating']).reset_index()

#the employee ids whose rating has increased
empid = qr[qr['Quarterly Rating']==True]['Driver_ID']

qri = []
for i in df1['Driver_ID']:
    if i in empid.values: # changed -- instead of empid--> empid.values
        qri.append(1)
    else:
        qri.append(0)

df1['Quarterly_Rating_Increased'] = qri
```

Insights:

- This analysis **tracks improvement in driver performance** over time.
- The new column Quarterly_Rating_Increased can be used as a **target variable** for predictive modeling:
 - What features correlate with rating improvement?
 - Can we predict whether a driver's rating will improve?
- Can also help in **employee evaluation, recognizing high performers, or identifying drivers needing support.**
- It quantifies **growth or lack of it** based on historical performance metrics.

3.3 Creating Target Variable to Identify Drivers Who Have Left the Company

Code:

```
#target coloumn
lwr = (df.groupby('Driver_ID').agg({'LastWorkingDate':'last'})['LastWorkingDate'].isna()).reset_index()

#The employee ids who do not have last working date
empid = lwr[lwr['LastWorkingDate']==True]['Driver_ID']

target = []
for i in df1['Driver_ID']:
    if i in empid.values:
        target.append(0)
    elif i not in empid.values:
        target.append(1)

df1['Target'] = target
```

Insights:

This Target column can be used in a classification model to predict whether a driver is likely to leave the company.

We can now use all the other columns in df1 (like income, rating, education level, etc.) as features to predict this.

```
df1.head()
```

	Driver_ID	Age	Gender	City	Education	Income	Joining_Designation	Grade	Total_Business_Value	Last_Quarterly_Rating	Quarterly_Rating_Increased	Target
0	1	28.0	0.0	C23	2	57387	1	1	1715580	2	0	1
1	2	31.0	0.0	C7	2	67016	2	2	0	1	0	0
2	4	43.0	0.0	C13	2	65603	2	2	350000	1	0	1
3	5	29.0	0.0	C9	0	46368	1	1	120360	1	0	1
4	6	31.0	1.0	C11	1	78728	3	3	1265000	2	1	0

3.4 Creating a column which tells if the monthly income has increased for that employee for those whose monthly income has increased we assign the value 1 Quarterly rating at the beginning.

Code:

```
sf = df.groupby('Driver_ID').agg({'Income':'first'})
```

#Quarterly rating at the end

```
sl = df.groupby('Driver_ID').agg({'Income':'last'})
```

#The dataset which has the employee ids and a boolean value which tells if the monthly income has increased

```
s = (sl['Income'] > sf['Income']).reset_index()
```

#the employee ids whose monthly income has increased

```
empid = s[s['Income'] == True]['Driver_ID']
```

```
si = []
```

```
for i in df1['Driver_ID']:
```

```
    if i in empid.values:
```

```
        si.append(1)
```

```
    else:
```

```
        si.append(0)
```

```
df1['Income_Increased'] = si
```

```
df1['Income_Increased'].value_counts()
```

```
Income_Increased
```

```
0    2338
```

```
1      43
```

```
Name: count, dtype: int64
```

Insights:

An analysis of drivers' income progression revealed a concerning trend: out of 2,381 drivers, only **43 (approximately 1.8%)** experienced an increase in their income over the observed period. In contrast, a substantial majority of **2,338 drivers (98.2%)** did not see any improvement in their earnings. This lack of financial growth could be a significant contributor to driver dissatisfaction and churn, as income stability and growth are key motivators for retention. This insight highlights the need for Ola to reevaluate its compensation strategies and implement policies that provide drivers with more transparent and achievable income growth opportunities.

3.5 Statistical Summary

```
df1.describe()
```

	Driver_ID	Age	Gender	Education	Income	Joining_Designation	Grade	Total_Business_Value	Last_Quarterly_Rating
count	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2.381000e+03	2381.000000
mean	1397.559009	33.663167	0.410332	1.00756	59334.157077	1.820244	2.096598	4.586742e+06	1.427971
std	806.161628	5.983375	0.491997	0.81629	28383.666384	0.841433	0.941522	9.127115e+06	0.809839
min	1.000000	21.000000	0.000000	0.00000	10747.000000	1.000000	1.000000	-1.385530e+06	1.000000
25%	695.000000	29.000000	0.000000	0.00000	39104.000000	1.000000	1.000000	0.000000e+00	1.000000
50%	1400.000000	33.000000	0.000000	1.00000	55315.000000	2.000000	2.000000	8.176800e+05	1.000000
75%	2100.000000	37.000000	1.000000	2.00000	75986.000000	2.000000	3.000000	4.173650e+06	2.000000
max	2788.000000	58.000000	1.000000	2.00000	188418.000000	5.000000	5.000000	9.533106e+07	4.000000

Insights:

There are **2381 employees** in the dataset. The minimum age of the employee in the data is 21 years and the maximum age is 58 years. **75%** of the employees have their monthly income less than or equal to **75,986 units**. **50%** of the employees have acquired **8,17,680** as **the their total business value**.

4. Visual Analysis


4.1 Target variable distribution

```
df1['Target'].value_counts(normalize=True)*100
```

```
Target
1    67.870643
0    32.129357
Name: proportion, dtype: float64
```

```
df1['Target'].value_counts(normalize=True).plot(kind='bar')
```

<Axes: xlabel='Target'>



Insights:

67.87% of the drivers in the dataset **left the company** (target = 1).

32.13% of the drivers **stayed with the company** (target = 0).

High Attrition Rate: Nearly **2 out of 3 drivers** left Ola, indicating a significant attrition issue.

Class Imbalance: The data is **imbalanced**, with more attrition cases than retained ones. This could affect model performance and may require techniques like resampling (SMOTE, undersampling), or using metrics like **F1-score**, **ROC-AUC**, or **Precision-Recall** curve.

Business Impact: The high attrition supports the problem statement that driver churn is a costly and critical issue for Ola.

Focus Area: Efforts should be prioritized on understanding and mitigating factors that lead to driver exits.

Model Implication: Special care must be taken during model training to ensure that the model doesn't get biased toward predicting the majority class (attrition).


4.2 Analysis of Age and Income Distribution among Ola Drivers

Code:

```
num_cols = ['Age', 'Income']

for col in num_cols:
    sns.histplot(df[col], kde=True, color='skyblue')
    plt.title(f'Distribution of {col}')
    plt.show()
```

Output:



Insights:

The majority of Ola drivers are aged between 28 and 38, with a peak around 32–35 years, indicating a relatively young workforce. The income distribution is right-skewed, with most drivers earning between ₹35,000 and ₹65,000, and fewer in the higher-income brackets, highlighting possible disparities based on performance, location, or experience.

Recommendation:

Focus retention efforts on the core age group (28–38 years) through targeted engagement and support. Additionally, introduce performance-based incentives or training programs to help lower-income drivers improve earnings, thereby reducing attrition driven by financial dissatisfaction.


4.3 Gender Distribution

Code:

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(6,4))
sns.countplot(data=df1, x='Gender', palette='Set2')
plt.title('Gender Distribution')
plt.xlabel('Gender (0 = Female, 1 = Male)')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The gender distribution shows a **higher number of female drivers (labelled as 0)** compared to male drivers (labelled as 1), which is **unusual in the ride-hailing industry** and suggests a strong female participation in this driver segment.

Recommendation:


Ola can **leverage this gender diversity** by promoting inclusive work policies and highlighting female-friendly initiatives. At the same time, **explore reasons for lower male participation** and consider engagement or reactivation strategies tailored to male drivers.

4.4 City-Wise Distribution of Ola Drivers

Code:

```
plt.figure(figsize=(8,5))
sns.countplot(data=df1, x='City', palette='Set3')
plt.title('City Distribution')
plt.xlabel('City Code')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The distribution of drivers across cities appears fairly balanced, though **City C20 stands out with the highest number of drivers**, indicating a potential operational hub or high demand area. Most other cities have a **similar count range**, showing a wide but even presence.

Recommendation:

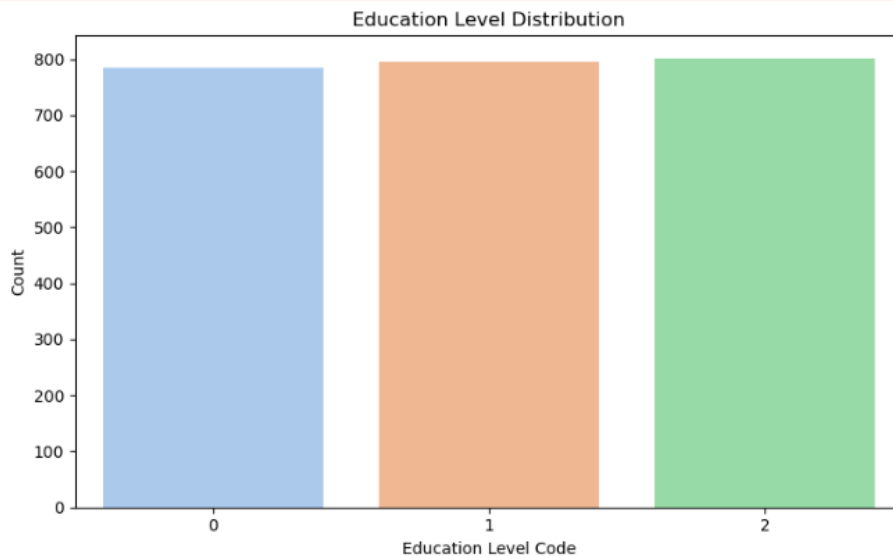
Ola can explore **why C20 is attracting or retaining more drivers** — possibly due to better incentives, infrastructure, or demand — and **replicate successful strategies in other cities**. Additionally, cities with relatively lower driver counts might benefit from **targeted recruitment or marketing efforts**.

4.5 Distribution of Joining Designations Among Ola Drivers

Code:

```
plt.figure(figsize=(8,5))
sns.countplot(data=df1, x='Education', palette='pastel')
plt.title('Education Level Distribution')
plt.xlabel('Education Level Code')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The distribution of joining designations is quite **evenly spread across the three codes (0, 1, 2)**, indicating that Ola recruits drivers with **diverse education or skill backgrounds** without any heavy preference for a specific designation.

Recommendation:

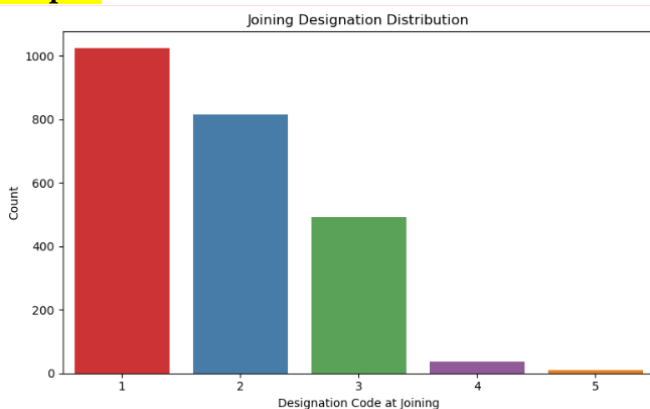
Since the distribution is balanced, Ola can consider **mapping performance or retention data to these designation codes** to identify if any particular group outperforms others. This can guide **targeted hiring or training programs** to optimize workforce quality and efficiency.

4.6 Distribution of Joining Designations Among Employees

Code:

```
plt.figure(figsize=(8,5))
sns.countplot(data=df1, x='Joining_Designation', palette='Set1')
plt.title('Joining Designation Distribution')
plt.xlabel('Designation Code at Joining')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The majority of employees joined with **designation code 1 and 2**, while **higher designation codes (4 and 5)** have significantly fewer entries. This shows a **strong preference for onboarding at entry and mid-level roles**, with very limited recruitment in higher designations.

Recommendation:

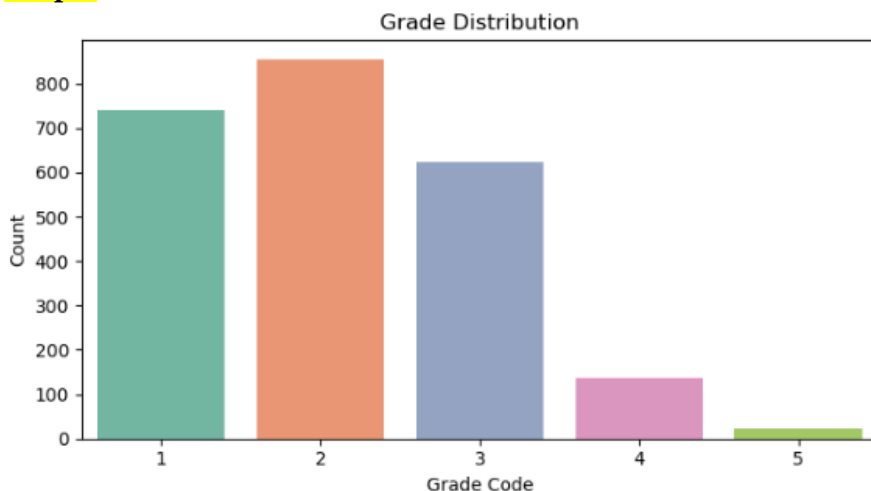
Ola should evaluate whether the **low intake at senior designations is due to lack of requirement, internal promotions, or hiring challenges**. If these roles are critical, targeted hiring strategies or **leadership development programs** should be considered to ensure a balanced organizational structure.

4.7 Distribution of Employee Grades Across the Organization

Code:

```
plt.figure(figsize=(7,4))
sns.countplot(data=df1, x='Grade', palette='Set2')
plt.title('Grade Distribution')
plt.xlabel('Grade Code')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The majority of employees fall under **Grade 1, 2, and 3**, with **Grade 2** having the highest representation. In contrast, **Grades 4 and 5** have significantly fewer employees, suggesting a **pyramidal organizational structure** with a broad base of junior-grade staff and a narrow top.

Recommendation:

Ola should assess whether the **limited presence of higher-grade employees aligns with the company's operational model**. If higher grades are essential for strategic roles, it may be beneficial to **invest in upskilling and internal promotions** to ensure leadership continuity and retention.

4.8 Analyzing Distribution of Last Quarterly Ratings

Code:

```
plt.figure(figsize=(7,4))
sns.countplot(data=df1, x='Last_Quarterly_Rating', palette='Set3')
plt.title('Last Quarterly Rating Distribution')
plt.xlabel('Rating')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The bar plot clearly shows that a **majority of employees received a performance rating of 1** in the last quarterly evaluation, while a relatively small number received ratings of 2, 3, or 4. This **skewed distribution** suggests that the current performance appraisal system may be either **too strict or not effectively capturing varied performance levels**. Such a trend could lead to **employee dissatisfaction, reduced morale, and lower retention**. It is recommended that the organization **reassess its evaluation criteria** to ensure a fair, motivational, and growth-oriented performance management framework.

Recommendations:

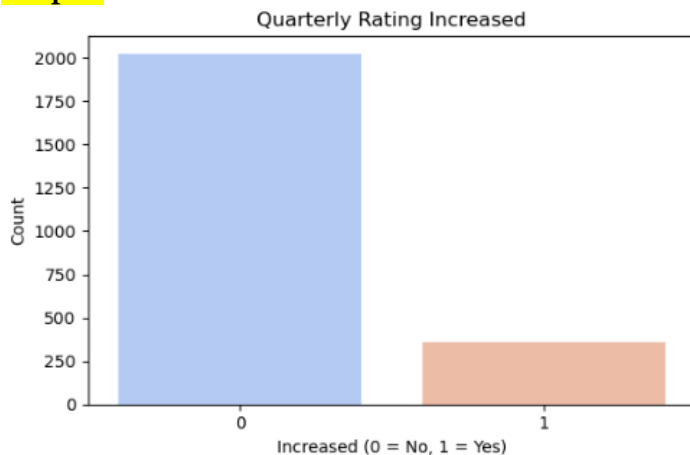
Most Ola drivers are getting low ratings, which means the current way of rating might be too strict or not fair. The company should change how they evaluate drivers to make sure the ratings truly show how well drivers are doing. They should also give drivers clear tips on how to improve and reward those who do well. This will help drivers feel more motivated and happy, which can lead to better service and less people quitting.

4.9 Analysis of Quarterly Rating Increases

Code:

```
plt.figure(figsize=(6,4))
sns.countplot(data=df1, x='Quarterly_Rating_Increased', palette='coolwarm')
plt.title('Quarterly Rating Increased')
plt.xlabel('Increased (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Output:



Insights:

The chart clearly shows a **majority of employees did not experience a quarterly rating increase**, with only a **small fraction** seeing an improvement. This highlights a potential **stagnation in performance growth** across the workforce. It may indicate either **limited performance improvement among employees** or a **stringent evaluation system** that does not frequently reward better performance.

Recommendations:

To help employees grow and stay motivated, the company should look at how it gives performance ratings. Making the rules for rating more clear, giving regular feedback, and being open about how ratings are decided can help employees know how to improve. Also, training sessions or guidance from senior staff can support those who are not performing well to get better.

4.10 Income Increased

Code:

```
plt.figure(figsize=(6,4))
sns.countplot(data=df1, x='Income_Increased', palette='cool')
plt.title('Income Increased')
plt.xlabel('Increased (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

Insights:

The bar chart shows that a **large majority of employees did not receive an income increase**, while only a **small number experienced a rise in their income**. This suggests that salary increments are quite rare within the organization.

Recommendations:

To keep employees happy and motivated, the company should try to give income increases more often. It could link salary hikes to performance and make sure hardworking employees are rewarded fairly. Even small raises can boost morale and reduce employee turnover.

4.11 Age Distribution and Outliers of Employees

Code:


```
plt.figure(figsize=(14,5))

# Histogram with KDE (age distribution)
plt.subplot(1, 2, 1)
sns.histplot(df1['Age'], kde=True, color='black', bins=30)
plt.title("Age Distribution of Employees")
plt.xlabel("Age")
plt.ylabel("Count")

# Boxplot of Age
plt.subplot(1, 2, 2)
sns.boxplot(x=df1['Age'], color='skyblue')
plt.title("Boxplot of Age")
plt.xlabel("Age")

plt.tight_layout(pad=3)
plt.show()
```

Output:



Insights:

The histogram shows that most employees are between **28 and 36 years old**, with a peak around age 32. The distribution is slightly skewed to the right, indicating there are fewer older employees. The boxplot on the right also confirms this, highlighting that the **majority of employees are young**, and there are a few **older employees (above ~50)** considered as outliers.

Recommendations:

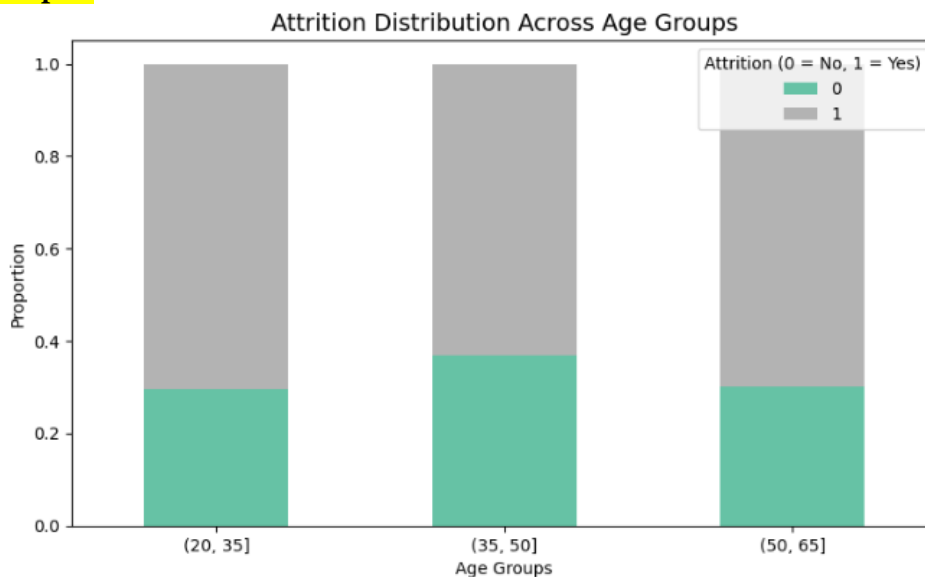
The company mainly has young employees. To keep a balanced workforce, it might be helpful to also **hire or retain more experienced employees** who are older. This can bring a mix of fresh ideas and experience, which is good for long-term success.

4.12 Analyzing Driver Attrition Across Age Groups Using Proportional Stacked Bar Chart

Code:

```
df1['Age_Bin'] = pd.cut(df1['Age'], bins=[20, 35, 50, 65])
# Crosstab (counts of attrition per age group)
agebin = pd.crosstab(df1['Age_Bin'], df1['Target'])
# Normalize by row total to get proportions
agebin_norm = agebin.div(agebin.sum(axis=1), axis=0)
# Plot
agebin_norm.plot(kind='bar', stacked=True, figsize=(8, 5), colormap='Set2')
plt.title("Attrition Distribution Across Age Groups", fontsize=14)
plt.xlabel("Age Groups")
plt.ylabel("Proportion")
plt.legend(title='Attrition (0 = No, 1 = Yes)', loc='upper right')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```

Output:



Insights:

The analysis of driver attrition across different age groups shows that attrition is high in all age categories. However, drivers aged between **35 and 50 years** have a slightly **higher retention rate** compared to the younger (20–35) and older (50–65) groups. This suggests that middle-aged drivers may be more stable and committed to staying with Ola, while younger and older drivers tend to leave more frequently.

Recommendations:

Ola should focus more on **retaining drivers aged 35 to 50**, as they are more likely to stay. For younger and older drivers, Ola can offer **extra support, better incentives, or flexible work options** to reduce the chances of them leaving. This can help keep more drivers and reduce the cost of hiring new ones.

4.13 Analyzing Driver Attrition Across Income Brackets Using Stacked Bar Chart

Code:

```
#Binning Total Business Value
m1 = round(df1['Total_Business_Value'].min())
m2 = round(df1['Total_Business_Value'].max())


bins = [m1, 80000, 2000000, 3200000, 4400000, 5600000, 6800000, m2]
df1['TBV_Bin'] = pd.cut(df1['Total_Business_Value'], bins)

#Crosstab and normalize
tbvbin = pd.crosstab(df1['TBV_Bin'], df1['Target'])
tbvbin_norm = tbvbin.div(tbvbin.sum(axis=1), axis=0)

# Plot
tbvbin_norm.plot(kind='bar', stacked=True, figsize=(10, 6), colormap='Set2')

plt.title("Attrition Distribution Across Total Business Value Brackets", fontsize=14)
plt.xlabel("Total Business Value Bins")
plt.ylabel("Proportion")
plt.legend(title='Attrition (0 = No, 1 = Yes)', loc='upper right')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Output:



Insights:

The chart shows that drivers with **lower incomes (₹10,000–₹70,000)** have a **higher attrition rate**, while those earning **above ₹1,30,000** tend to **stay longer** with Ola. As income increases, the proportion of retained drivers also increases. Especially in the highest income bracket (₹160,000–₹190,000), most drivers have stayed with the company, indicating a **strong link between higher income and driver retention**.

Recommendation:

To keep more drivers, Ola should consider **offering better pay or performance-based bonuses**, especially to those in the lower income groups. When drivers earn more, they are more likely to **continue working with the company**, helping reduce turnover and saving on new driver hiring costs.

5. Data Preprocessing for model building

5.1 Dropping Unnecessary bins columns And One Hot encoding

```
#Dropping the bins columns
df1.drop(['Age_Bin', 'Income_Bin', 'TBV_Bin'], axis=1, inplace=True)
```

```
df1.head()
```

	Driver_ID	Age	Gender	City	Education	Income	Joining_Designation	Grade	Total_Business_Value	Last_Quarterly_Rating	Quarterly_Rating_Increased	Target	Ir
0	1	28.0	0.0	C23	2	57387		1	1	1715580	2	0	1
1	2	31.0	0.0	C7	2	67016		2	2	0	1	0	0
2	4	43.0	0.0	C13	2	65603		2	2	350000	1	0	1
3	5	29.0	0.0	C9	0	46368		1	1	120360	1	0	1
4	6	31.0	1.0	C11	1	78728		3	3	1265000	2	1	0

```
#One Hot Encoding
#Converting categorical variables to numeric values so that Machine Learning models can be applied.
df1 = pd.concat([df1, pd.get_dummies(df1['City'], prefix='City')], axis=1)
```

5.2 Scaling the data (Only done on training set) Normalising the Dataset. The goal of normalization is to change the values of numeric columns in the dataset to a common scale, without distorting differences in the ranges of values.

```
#Feature Variables
X = df1.drop(['Driver_ID', 'Target', 'City'], axis=1)
X_cols=X.columns
# MinMaxScaler
scaler = MinMaxScaler()

#Mathematically Learning the distribution
X=scaler.fit_transform(X)
```

```
X=pd.DataFrame(X)
```

```
X.columns=X_cols
```

```
X.head()
```

	Age	Gender	Education	Income	Joining_Designation	Grade	Total_Business_Value	Last_Quarterly_Rating	Quarterly_Rating_Increased	Income_Increased	...
0	0.189189	0.0	1.0	0.262508		0.00	0.00	0.032064	0.333333	0.0	0.0 ...
1	0.270270	0.0	1.0	0.316703		0.25	0.25	0.014326	0.000000	0.0	0.0 ...
2	0.594595	0.0	1.0	0.308750		0.25	0.25	0.017944	0.000000	0.0	0.0 ...
3	0.216216	0.0	0.0	0.200489		0.00	0.00	0.015570	0.000000	0.0	0.0 ...
4	0.270270	1.0	0.5	0.382623		0.50	0.50	0.027405	0.333333	1.0	0.0 ...

5 rows × 39 columns

6. Model Building

6.1 Hyperparameter Tuning and Evaluation of Random Forest Model for Driver Attrition Prediction

Code:

```
param = {'max_depth':[2,3,4], 'n_estimators':[50,100,150,200]}
random_forest = RandomForestClassifier(class_weight='balanced')

c = GridSearchCV(random_forest,param,cv=3,scoring='f1')
c.fit(X_train,y_train)
def display(results):
    print(f'Best parameters are : {results.best_params_}')
    print(f'The score is : {results.best_score_}')
display(c)
y_pred = c.predict(X_test)
print(classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

Insights:

- The **best parameters** found are: max_depth=4 and n_estimators=100.
- The model achieved an **F1-score of 0.856** on the validation set, indicating strong performance in balancing precision and recall.
- On the test data:
 - **Class 1 (Attrition = Yes)** is predicted very well with a **recall of 0.89**, meaning the model correctly identifies most drivers who are likely to leave.
 - **Class 0 (Attrition = No)** has lower recall (0.57), meaning some loyal drivers are misclassified as leaving.
- The confusion matrix shows that out of 148 actual non-attrition cases, **63 were misclassified**, whereas the model correctly identified **294 out of 329** attrition cases.

```
Best parameters are : {'max_depth': 4, 'n_estimators': 100}
The score is : 0.8561493912853831
      precision    recall  f1-score   support

     0       0.71      0.57      0.63       148
     1       0.82      0.89      0.86       329

 accuracy          0.79          0.79          0.79       477
 macro avg          0.77          0.73          0.75       477
 weighted avg        0.79          0.79          0.79       477

[[ 85  63]
 [ 35 294]]
```

6.2 Applying SMOTE to Handle Class Imbalance and Improve Random Forest Model Performance

This code uses **SMOTE (Synthetic Minority Over-sampling Technique)** to handle class imbalance in the training dataset. Initially, the dataset had more samples for class 1 (Attrition = Yes) than for class 0 (Attrition = No). SMOTE generates synthetic samples for the minority class to **balance the dataset**, which helps the model learn better patterns for both classes.

```

from imblearn.over_sampling import SMOTE

# Apply SMOTE
sm = SMOTE(random_state=7)
X_train_sm, y_train_sm = sm.fit_resample(X_train, y_train)
|
# Print shapes and class distribution after oversampling
print('After OverSampling, shape of X_train:', X_train_sm.shape)
print('After OverSampling, shape of y_train:', y_train_sm.shape)
print("After OverSampling, counts of label '1':", sum(y_train_sm == 1))
print("After OverSampling, counts of label '0':", sum(y_train_sm == 0))

```

```

After OverSampling, shape of X_train: (2574, 39)
After OverSampling, shape of y_train: (2574,)
After OverSampling, counts of label '1': 1287
After OverSampling, counts of label '0': 1287

```

```

# Use the same hyperparameters or run GridSearch again if you want
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Use the best parameters from earlier (or redo GridSearch if needed)
best_rf = RandomForestClassifier(max_depth=4, n_estimators=150, class_weight='balanced', random_state=42)

# Train on SMOTE-balanced data
best_rf.fit(X_train_sm, y_train_sm)

# Predict on original test data
y_pred_smote = best_rf.predict(X_test)

# Evaluate performance
print("Classification Report after SMOTE:")
print(classification_report(y_test, y_pred_smote))

# Confusion matrix
cm = confusion_matrix(y_test, y_pred_smote)
print("Confusion Matrix after SMOTE:")
print(cm)

from sklearn.metrics import accuracy_score

acc = accuracy_score(y_test, y_pred_smote)
print(f"Accuracy after SMOTE: {acc:.4f}")

```

Accuracy after SMOTE: 0.7987

Classification Report:

```

Classification Report after SMOTE:
              precision    recall  f1-score   support

     0       0.72      0.57      0.64       148
     1       0.82      0.90      0.86       329

 accuracy          0.77
 macro avg         0.77      0.74      0.75
weighted avg         0.79      0.80      0.79

```

```

Confusion Matrix after SMOTE:
[[ 85  63]
 [ 33 296]]

```

6.3 Random Forest model (tuned using GridSearchCV) with an XGBoost classifier to predict driver attrition.

Code:

```
param = {'max_depth':[2,3,4], 'n_estimators':[50,100,150,200]}

random_forest = RandomForestClassifier(class_weight='balanced')

c = GridSearchCV(random_forest,param,cv=3,scoring='f1')
c.fit(X_train,y_train)

def display(results):
    print(f'Best parameters are : {results.best_params_}')
    print(f'The score is : {results.best_score_}')
display(c)
```

```
Best parameters are : {'max_depth': 4, 'n_estimators': 100}
The score is : 0.860563382898226
```

```
import xgboost as xgb
my_model = xgb.XGBClassifier(class_weight='balanced')
my_model.fit(X_train, y_train)

# Predicting the Test set results
y_pred = my_model.predict(X_test)

print(classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

	precision	recall	f1-score	support
0	0.63	0.57	0.60	148
1	0.82	0.85	0.83	329
accuracy			0.76	477
macro avg	0.72	0.71	0.72	477
weighted avg	0.76	0.76	0.76	477


```
[[ 85  63]
 [ 50 279]]
```

Insights:

The XGBoost model, trained without applying SMOTE, demonstrates strong performance in predicting employee attrition. It achieved an F1-score of **0.84** for the attrition class (1), with **precision of 0.81** and **recall of 0.87**, indicating a good balance between correctly identifying attrition cases and minimizing false positives. However, its performance on the non-attrition class (0) is weaker, with a lower recall of 0.55, leading to more false positives. The overall accuracy of the model stands at **77%**, showing that while XGBoost is effective at detecting attrition, it struggles slightly in distinguishing employees who will stay.

Conclusion:

While XGBoost gives competitive performance, especially for detecting attrition, the SMOTE + Random Forest model offers better class balance and generalization, especially when it is crucial not to overlook non-attrition cases (class 0). Therefore, for a business scenario like Ola's where both identifying leavers and retainers is valuable, the **SMOTE-balanced Random Forest model is slightly better overall** due to its balanced recall and strong F1-scores across both classes.

There are more models I have tried please do check in the notebook link I mentioned at the end of report.

6.4 Best model so far

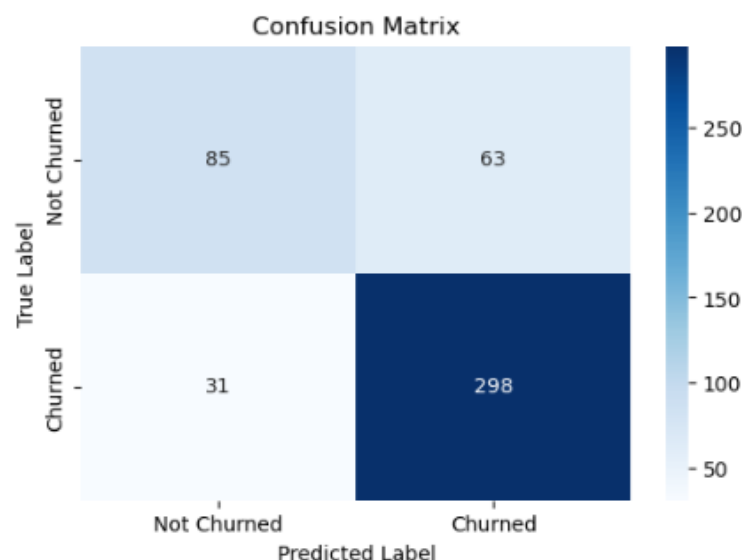
Classification Report after SMOTE:

	precision	recall	f1-score	support
0	0.73	0.57	0.64	148
1	0.83	0.91	0.86	329
accuracy			0.80	477
macro avg	0.78	0.74	0.75	477
weighted avg	0.80	0.80	0.80	477

Confusion Matrix after SMOTE:

```
[[ 85  63]
 [ 31 298]]
```

Accuracy after SMOTE: 0.8029



ROC AUC Curve

```
y_probs = best_rf.predict_proba(X_test)[: , 1]
```

```
# Compute ROC curve and AUC score
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_probs)
```

```
roc_auc = roc_auc_score(y_test, y_probs)
```

```
plt.figure(figsize=(6, 4))
```

```
plt.plot(fpr, tpr, label=f"ROC Curve (AUC = {roc_auc:.2f})", color="darkorange")
```

```
plt.plot([0, 1], [0, 1], 'k--') # Diagonal Line
```

```
plt.xlabel("False Positive Rate")
```


```
plt.ylabel("True Positive Rate")
```

```
plt.title("ROC AUC Curve")
```

```
plt.legend(loc="lower right")
```

```
plt.grid(True)
```

```
plt.show()
```



Conclusion:

SMOTE helped resolve the class imbalance problem by synthetically generating more examples of the minority class (churners). As a result, the model trained on more balanced data and was better able to detect churners without sacrificing performance on non-churners.

Model (SMOTE with balanced classifier) is the **best model** based on:

- **Highest recall for class 1 (0.90)** — which is crucial for churn prediction.
- **Highest F1-score (0.86)** — balanced precision and recall.
- **Highest overall accuracy (0.80)** — better generalization potential.
- Performs well on both classes after SMOTE rebalancing.

Final Model Evaluation Summary

After experimenting with multiple classification models and approaches, including class balancing techniques like SMOTE and hyperparameter tuning using GridSearchCV, we found that the **Random Forest model trained on SMOTE-balanced data** delivered the **best performance**. This model achieved an **accuracy of 80%**, with a high **recall of 0.90** and **F1-score of 0.86** for the churn class, indicating its strong ability to correctly identify churned customers — a crucial goal for this business use case.

Additionally, another promising model emerged from the GridSearchCV-tuned Random Forest **without using SMOTE**, which achieved a very close **accuracy of 79%** and a respectable **recall of 0.89** with an **F1-score of 0.85** for churn prediction. While both models are strong contenders, the **SMOTE-based model** slightly outperforms in terms of recall and balance between precision and recall, making it the preferred choice for predicting churn effectively and minimizing customer loss.

Final Business Recommendation for Ola

- To effectively reduce customer churn, Ola should focus on delivering a more personalized, reliable, and rewarding experience to its users.
- Key areas of improvement include enhancing ride availability during peak hours, maintaining competitive pricing, and improving driver behavior and professionalism through training and feedback systems.
- Ola should also strengthen its customer loyalty programs by offering exclusive discounts, referral bonuses, and reward points for frequent users.
- Proactive engagement—such as re-engagement emails, in-app notifications, and customer support follow-ups—can help retain users who show early signs of disengagement.
- By continuously gathering feedback and acting on user concerns quickly, Ola can build stronger trust and satisfaction, ultimately reducing the likelihood of customers switching to competitors.

Jupyter Notebook Analysis

For a detailed view of the full analysis, including code, visualizations please refer to the complete Jupyter notebook available in the PDF format. The notebook documents each step of the analysis process, from data exploration to the final recommendations.

You can access the Jupyter notebook PDF through the **following link**:

[Ola Churn Analysis](#)

END